

CAP 6315 - Social Networks and Big Data Analytics

Project Part #1

Jordan Winemiller - Z23962588

[Colab Project Link](#)

[GitHub Project Link](#)

This project will be using the [Twitch Social Network](#) datasets, and be examining the Brazilian Portuguese language stream for our network.

```
In [67]: # Setup and Imports
import random
import os
import zipfile

import matplotlib.pyplot as plt
import networkx as nx
import numpy as np
from pyspark.sql import (
    SparkSession,
    Window,
)
import pyspark.sql.functions as F
import pyspark.sql.types as T
import seaborn as sns

SEED = 1234
LANG = "PTBR"
PATH = os.getcwd()
DATA_PATH = os.path.join(
    PATH,
    "twitch",
    LANG,
)

random.seed(SEED)
np.random.seed(SEED)
```

Importing the data and creating PySpark session

Unzipping archive file

```
In [2]: archive_file = "twitch.zip"
with zipfile.ZipFile(os.path.join(PATH, archive_file), "r") as archive_obj:
    archive_obj.extractall(PATH)
```

Creating the Spark Session and Loading the Dataset

```
In [3]: spark = SparkSession.builder \
    .appName("ProjectPart-1") \
    .getOrCreate()
```

```
In [4]: df = spark.read.csv(
    path=os.path.join(DATA_PATH, f"musae_{LANG}_edges.csv"),
    header=True,
    inferSchema=True
)
df.head(5)
```

```
Out[4]: [Row(from=0, to=92),
Row(from=0, to=428),
Row(from=1, to=689),
Row(from=1, to=1147),
Row(from=1, to=1666)]
```

Creating Graph with NetworkX

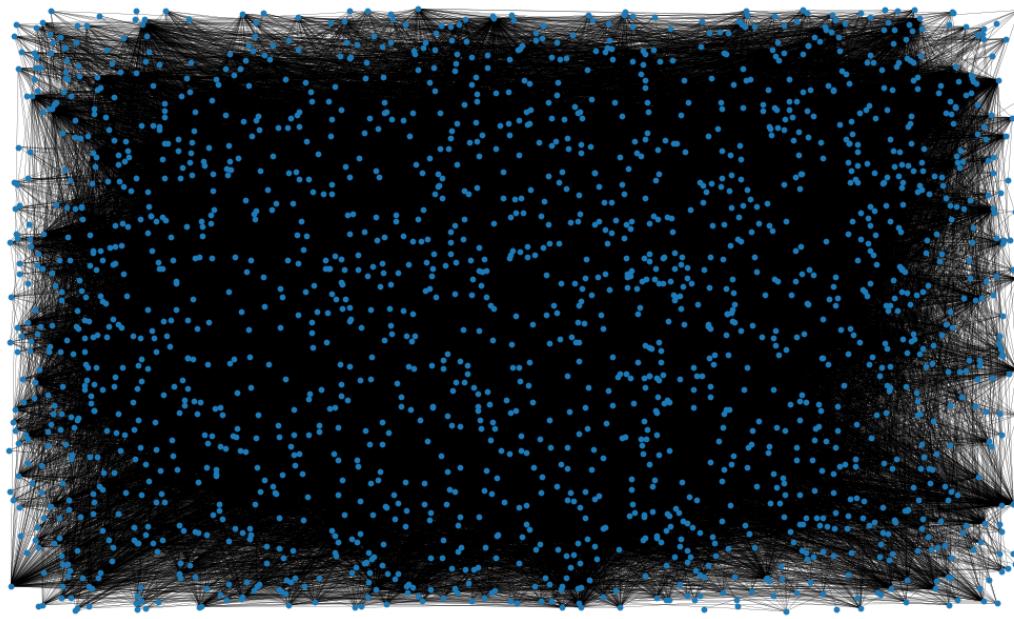
```
In [5]: G = nx.Graph()

edges = df.collect()
G.add_edges_from([
    (row["from"], row["to"])
    for row in edges
])
```

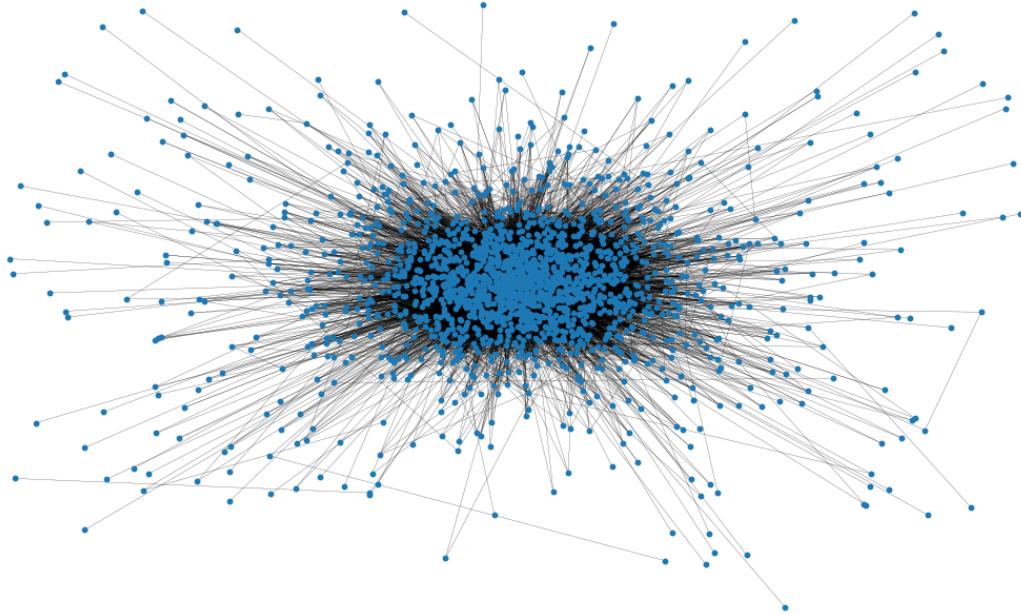
Graph Visualization

```
In [6]: plot_options = {
    "node_size": 10,
    "with_labels": False,
    "width": 0.15,
}
```

```
In [7]: fig, ax = plt.subplots(figsize=(15, 9))
ax.axis("off")
nx.draw_networkx(
    G,
    pos=nx.random_layout(G),
    ax=ax,
    **plot_options)
```



```
In [8]: pos = nx.spring_layout(  
    G,  
    iterations=15,  
    seed=SEED,  
)  
fig, ax = plt.subplots(figsize=(15, 9))  
ax.axis("off")  
nx.draw_networkx(  
    G,  
    pos=pos,  
    ax=ax,  
    **plot_options)
```



Calculate the Number of Nodes in the Graph

```
In [9]: nn = G.number_of_nodes()
print(f"The Number of Nodes: {nn}")
```

The Number of Nodes: 1912

Calculate the Number of Edges in the Graph

```
In [10]: ne = G.number_of_edges()
print(f"The Number of Edges: {ne}")
```

The Number of Edges: 31299

Calculate the Network Diameter

Both methods are displayed

```
In [11]: min_path_len = dict(nx.all_pairs_shortest_path_length(G))
diameter = max(nx.eccentricity(G, sp=min_path_len).values())

print(f"Diameter with Eccentricity Method: {diameter}")
print(f"Diameter with NetworkX: {nx.diameter(G)}")
```

Diameter with Eccentricity Method: 7
Diameter with NetworkX: 7


Network Diameter Interpretation

The diameter of the network is **7**, which means that at most any user is **7** connections or friendships away from any other user in the network.

Histogram of Degree Distribution

```
In [12]: deg = [d for _, d in G.degree()]

plt.figure(figsize=(15, 9))
plt.hist(
    x=deg,
    bins=range(min(deg), max(deg) + 2),
    edgecolor="blue",
    align="left",
)
plt.title("Degree Distribution")
plt.xlabel("Degree")
plt.ylabel("Number of Nodes")
plt.grid(
    axis="y",
    linestyle="dashed",
    alpha=0.7,
)
plt.show()
```



Histogram Observations

The distribution is very skewed which would follow a power-law distribution. So, a lot of users have a small amount of friends and towards the right side of the graph a few users have a large amount of friends.

Average Degree of a Node

```
In [13]: avg_deg = np.mean([d for _, d in G.degree()])
print(f"Average Degree of a Node: {avg_deg}")
```

Average Degree of a Node: 32.73953974895397

Average Degree in Dataset Context

On average, a user has almost **33** friends in the network.

Calculate the Average Shortest Path Length Between Any Two Nodes

```
In [14]: avg_path_len = [
    np.mean(list(mpl.values()))
    for mpl in min_path_len.values()
]
print(f"Average Shortest Path Length: {np.mean(avg_path_len)}")
```

Average Shortest Path Length: 2.5310546909192766

Average Shortest Path Dataset Context

On average, a user is **3** friendships away from another user in the network.

Plot the Distribution of the Shortest Path Lengths

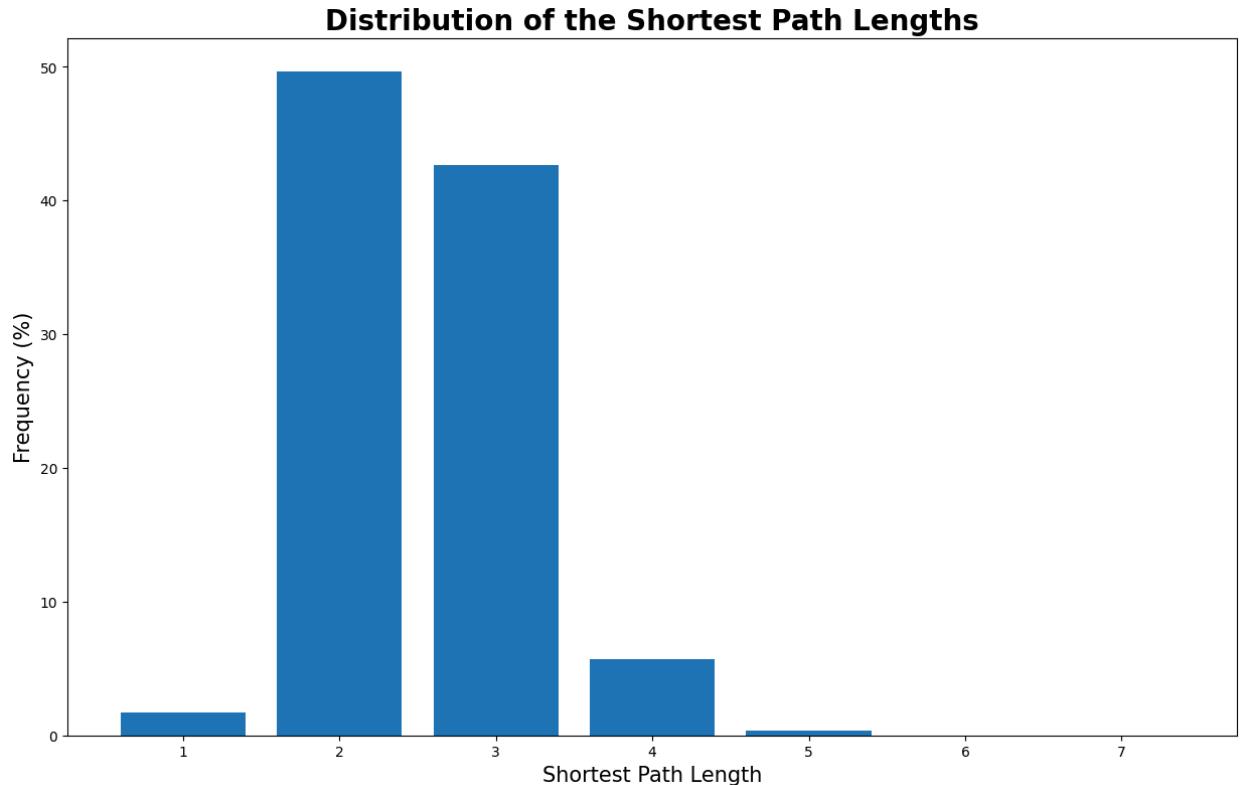
```
In [15]: path_len = np.zeros(diameter + 1, dtype=int)

for pls in min_path_len.values():
    pl, cnt = np.unique(list(pls.values()), return_counts=True)
    path_len[pl] += cnt

freq_percent = 100 * path_len[1:] / np.sum(path_len[1:])

fig, ax = plt.subplots(figsize=(15, 9))
ax.bar(np.arange(1, diameter + 1), height=freq_percent)
ax.set_title(
    label="Distribution of the Shortest Path Lengths",
    fontdict={
        "fontsize": 20,
        "fontweight": "bold",
    },
    loc="center",
)
ax.set_xlabel(
    xlabel="Shortest Path Length",
    fontdict={
        "fontsize": 15,
    },
)
ax.set_ylabel(
    ylabel="Frequency (%)",
    fontdict={
        "fontsize": 15,
    },
)
```

Out[15]: Text(0, 0.5, 'Frequency (%)')



Shortest Path Histogram Observations

The majority of the users are from 2 to 3 friends apart. These two shortest paths account for around 90% of how far a user is away from another user. So, the chance that any pair of users are 5, 6, or 7 (*7 being the network diameter*) friends apart is rare.

Calculate the Edge Density of this Graph

```
In [16]: print(f"Edge Density: {nx.density(G)}")
```

Edge Density: 0.017132150575067492

Conclusion on Histogram with Edge Density

We can conclude that the network is relatively sparse since only around 17% of the possible edges are present in the network.

Load the Node Attribute File and Display the First Few Rows of the Dataset

```
In [17]: att_df = spark.read.csv(
    path=os.path.join(DATA_PATH, f"musae_{LANG}_target.csv"),
    header=True,
    inferSchema=True
)

att_df.show(5)
```

```

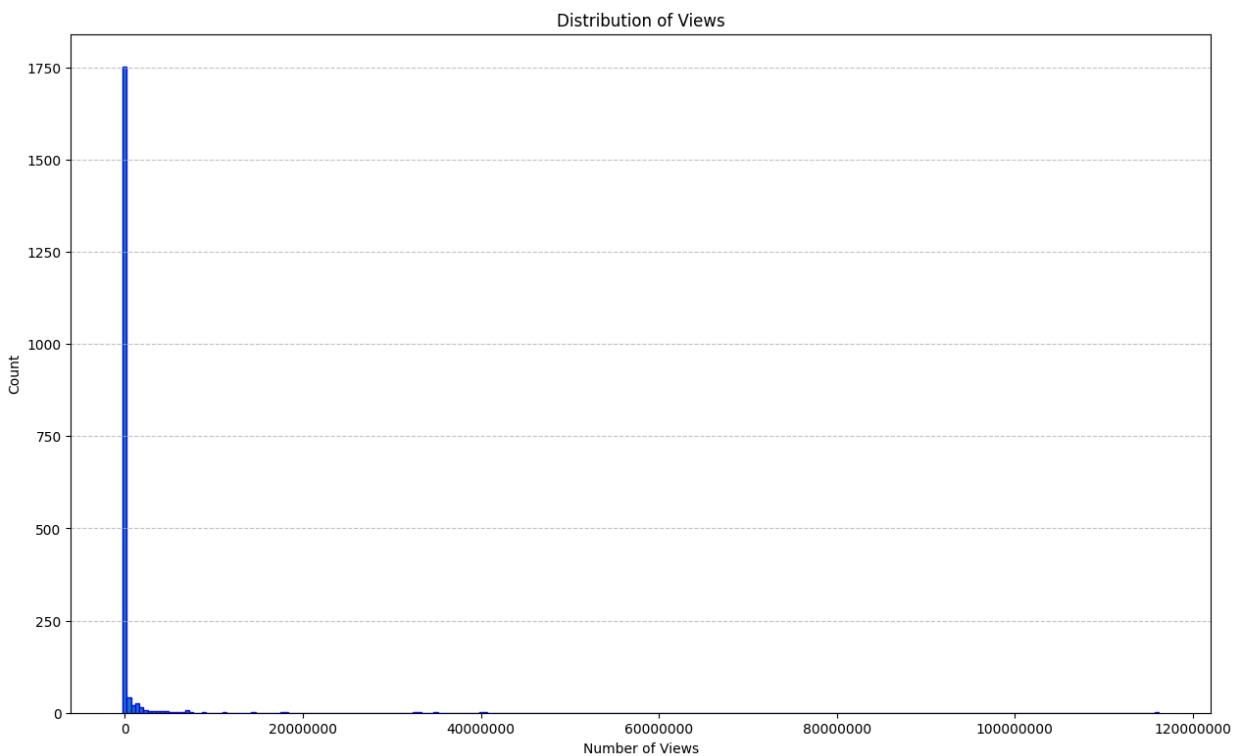
+-----+-----+-----+-----+
|      id|days|mature| views|partner|new_id|
+-----+-----+-----+-----+
| 44891403|1943| false|929459|   true| 1706|
| 61180621|1633| false| 11194|  false| 1273|
|145109685| 632| false|  2850|  false|  313|
|122121521| 906|  true|  3422|  false| 1570|
|189445819| 267| false|     71|  false|  800|
+-----+-----+-----+-----+
only showing top 5 rows

```

Plot a Histogram of the Views

```
In [18]: views = att_df.select("Views").rdd.map(lambda x: x[0]).collect()

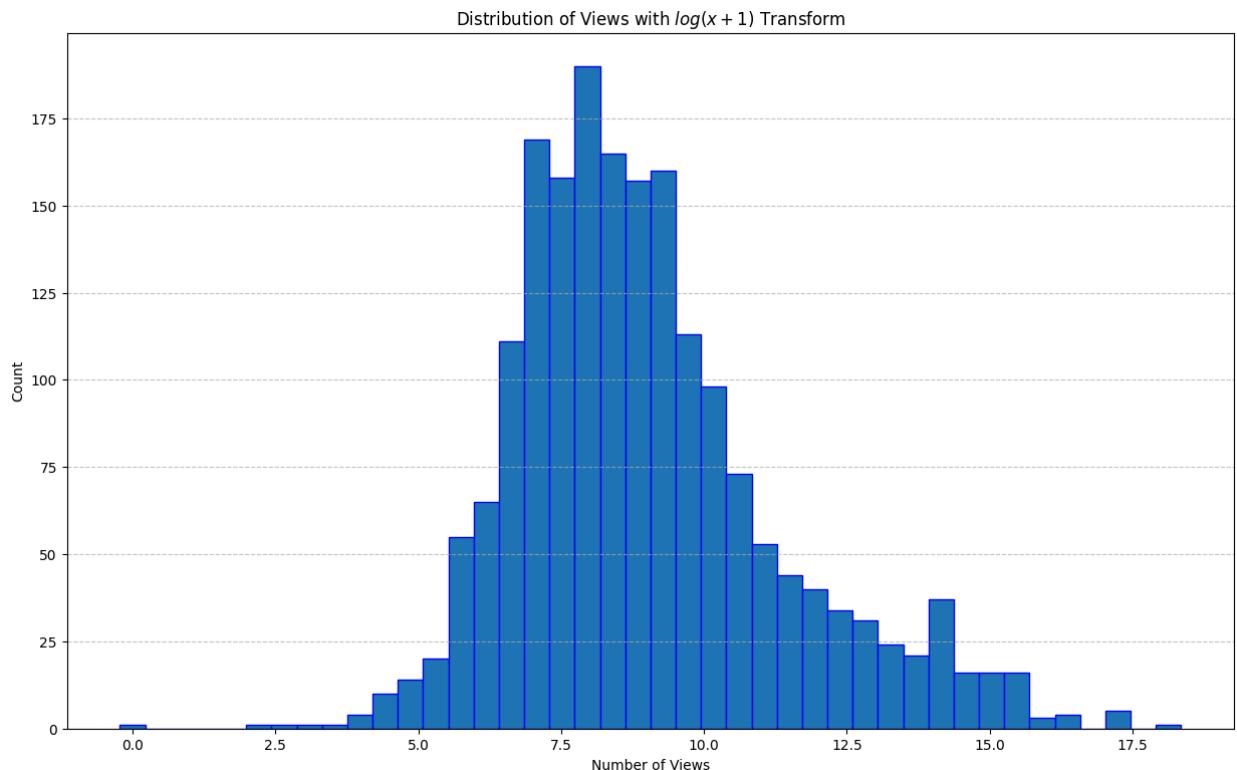
plt.figure(figsize=(15, 9))
plt.hist(
    x=views,
    bins=250,
    edgecolor="blue",
    align="left",
)
plt.title("Distribution of Views")
plt.xlabel("Number of Views")
plt.ticklabel_format(style="plain")
plt.ylabel("Count")
plt.grid(
    axis="y",
    linestyle="dashed",
    alpha=0.7,
)
plt.show()
```



Since there is a large skewness in the number of views; A Log Transformation $\log(x + 1)$ was preformed to normalize the data.

```
In [19]: log_views = np.log([v + 1 for v in views])

plt.figure(figsize=(15, 9))
plt.hist(
    x=log_views,
    bins="auto",
    edgecolor="blue",
    align="left",
)
plt.title("Distribution of Views with $log(x+1)$ Transform")
plt.xlabel("Number of Views")
plt.ticklabel_format(style="plain")
plt.ylabel("Count")
plt.grid(
    axis="y",
    linestyle="dashed",
    alpha=0.7,
)
plt.show()
```



Viewer Distribution Observations

There is a skewness to the distribution, and after the data was normalized with a log transformation we are able to see that there are also outliers in the dataset.

Create a Bar Plot

Showing the Number of Streamers who are Partners vs. Non-Partners

```
In [20]: att_df.groupBy("partner").count().show()

plt.bar(
    x = ["Partners", "Non-Partners"],
```


```

    height =att_df.groupBy("partner").count().rdd.map(lambda x: x[1]).collect(),
)
plt.title("Number of Streamers who are Partners vs. Non-Partners")
plt.ylabel("Number of Streamers")

plt.show()

```

partner	count
true	279
false	1633



Partnership Bar Chart Observations

- The proportion of streamers that are partnered make up around 85.4% of the dataset.

Create a Bar Plot

- Showing the Number of Streamers with Mature Content Enabled vs. Not Enabled

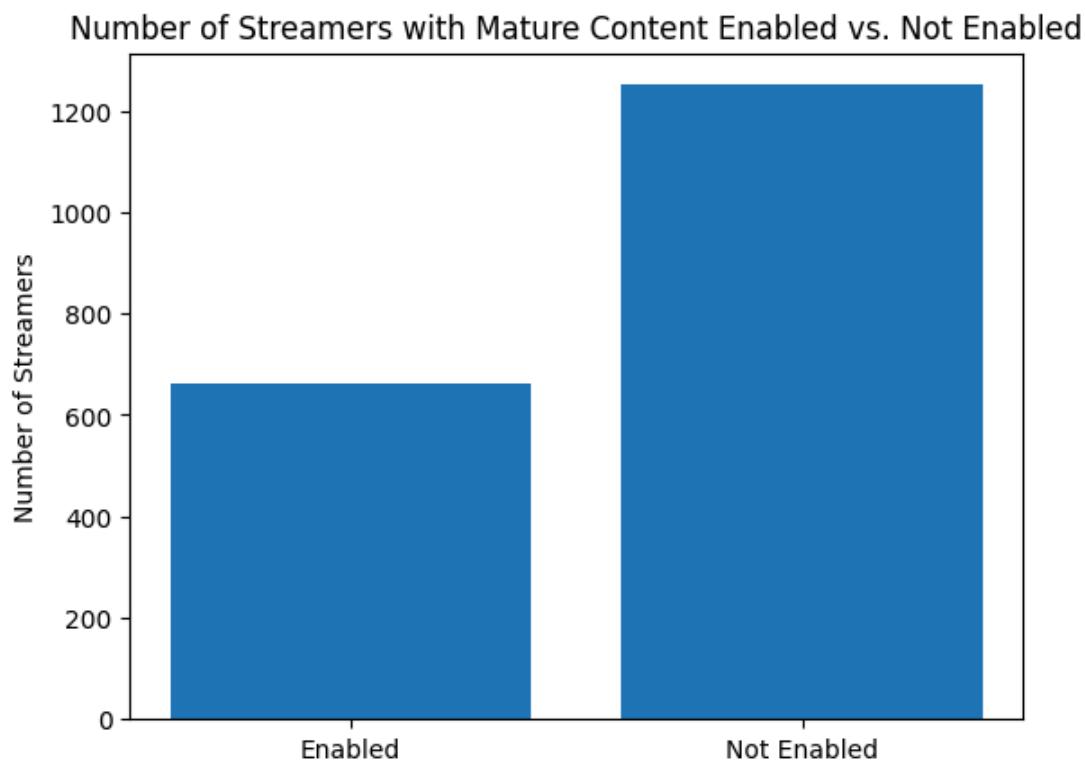
```

In [21]: att_df.groupBy("mature").count().show()

plt.bar(
    x=["Enabled", "Not Enabled"],
    height=att_df.groupBy("mature").count().rdd.map(lambda x: x[1]).collect(),
)
plt.title("Number of Streamers with Mature Content Enabled vs. Not Enabled")
plt.ylabel("Number of Streamers")
plt.show()

```

```
+-----+-----+
|mature|count|
+-----+-----+
|  true|   661|
| false| 1251|
+-----+-----+
```




Mature Content Bar Chart Observations

The majority of streamers do not have mature content enable with a little more 65.43% of streamers not enabling mature content.

Create a Boxplot Comparing Views by Partner Status

```
In [22]: sns.boxplot(
    x=att_df.select("Views").rdd.map(lambda x: x[0]).collect(),
    hue=att_df.select("Partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("Views By Parnter Status")
plt.xlabel("Views")
plt.ticklabel_format(axis="x", style="plain")
plt.show()
```

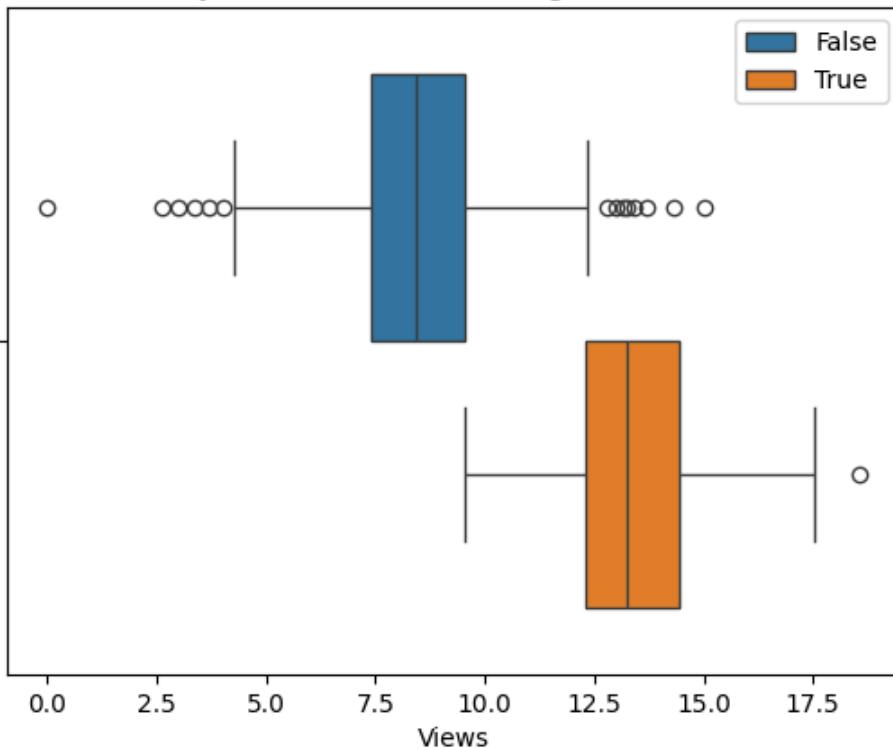
Views By Parnter Status



As before we applied the Log Transformation $\log(x + 1)$ on the number views to normalized the data.

```
In [23]: sns.boxplot(  
    x=log_views,  
    hue=att_df.select("Partner").rdd.map(lambda x: x[0]).collect(),  
)  
plt.title("Views By Parnter Status with $log(x+1)$ Transform")  
plt.xlabel("Views")  
plt.show()
```

Views By Parnter Status with $\log(x + 1)$ Transform



Views by Partnership Status Boxplot Observations

The initial boxplot shows that Partners do have higher views than Non-Partners by a large margin, this is reforced with boxplot that examines the log transformed data.

Compute for All Nodes

- Degree Centrality
- Betweenness Centrality
- Closeness Centrality
- Eigenvector Centrality
- PageRank

```
In [24]: degree = nx.degree_centrality(G)
betweenness = nx.betweenness_centrality(G)
closeness = nx.closeness_centrality(G)
eigenvector = nx.eigenvector_centrality(G)
pagerank = nx.pagerank(G)
```

```
In [25]: centrality_list = [
    {"degree centrality": degree},
    {"betweenness centrality": betweenness},
    {"closeness centrality": closeness},
    {"eigenvector centrality": eigenvector},
    {"pagerank": pagerank},
]

for centrality in centrality_list:
    print(f"Top 10 {list(centrality)[0].title()}")
    centrality_dict = centrality[list(centrality)[0]]
```

```
nodes = sorted(
    centrality_dict.items(),
    key=lambda x: x[1],
    reverse=True
)[:10]

for node in nodes:
    print(f"{node[0]}: {node[1]}")
print()
```

Top 10 Degree Centrality

127: 0.4013605442176871
1476: 0.3129251700680272
290: 0.3087388801674516
1297: 0.30716902145473574
467: 0.304552590266876
1660: 0.24856096284667714
67: 0.23757195185766616
1320: 0.217687074829932
1758: 0.20617477760334904
1259: 0.20146520146520147

Top 10 Betweenness Centrality

127: 0.0992613670254655
1476: 0.054894872658417665
1297: 0.05042095543765276
290: 0.05011201385679391
467: 0.04382473622801392
67: 0.034425706705748176
1660: 0.02624664025473586
1259: 0.023294665516722747
287: 0.022120764244835684
428: 0.021980008718637752

Top 10 Closeness Centrality

127: 0.6009433962264151
1297: 0.5684116597263533
467: 0.565050266114725
290: 0.5595900439238654
1476: 0.5594262295081968
67: 0.5435153583617748
1660: 0.5416666666666666
1593: 0.5295095594347464
1259: 0.5237051246916964
287: 0.5227024070021882

Top 10 Eigenvector Centrality

127: 0.17877831012617343
1297: 0.15866803272155214
467: 0.1572393891554249
290: 0.15150156437553064
1476: 0.1457607576802339
1660: 0.13375702630371897
67: 0.1311805297345441
1593: 0.12589762048531492
1320: 0.11858724919571517
1758: 0.11540890407339242

Top 10 Pagerank

127: 0.011868965496806964
1476: 0.008978180166428392
290: 0.008649101005769455
1297: 0.008500930100419827
467: 0.00842343460623767
1660: 0.00674011506372157
67: 0.006323889093141329
1320: 0.005883878568697245
1758: 0.005648521214050491
1259: 0.0056101915085292

Observations on Centrality Measures and PageRank

- Looking across the measures node 127 is at the top of each measure.
- Nodes 67, 290, 467, 1297, 1476 and 1660 appearing in the top 7 of each measure.
- Nodes 1320 and 1758 appears in *Degree* and *Eigenvector* centralities and *PageRank*, where node 1259 appears across all measures except for *Eigenvector* cetrality.
- Node 1593 is in the top 10 for only *Closeness* and *Eigenvector* centralities. Node 287 is only in the top 10 for *Betweenness* and *Closeness* centralities.
- Finally, node 428 is only in the top 10 *Betweenness* centrality across all the centralities and *PageRank*.


Create Boxplots Comparing Centralities and PageRang Values Between Partner Streamers and Non-Partner Streamers

Create the dataframe by joining the attributes data with centrality dictionary.

- Create PySpark dataframe from the centrality dictionary
- Join the dataframe to the attributes data to established partnership
- Visualize the boxplot for each centrality

```
In [26]: cent_df = spark.createDataFrame(
    degree.items(), ["node", "centrality"])
cent_df = cent_df.join(
    other=att_df.select("new_id", "Partner"),
    on=[cent_df.node == att_df.new_id],
    how="inner",
)
plt.figure(figsize=(15, 9))
sns.boxplot(
    x=cent_df.select("centrality").rdd.map(lambda x: x[0]).collect(),
    hue=cent_df.select("partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("Degree Centrality By Partnership Status")
plt.xlabel("Degree Centrality Values")
plt.show()
```

Degree Centrality By Partnership Status




Degree Centrality Boxplot

The partnered streamers have a lot more direct friendships than the non-partnered streamers. From the boxplot for *degree centrality* the largest > outlier in the non-partnered streamers would not be an outlier among the partnered streamers, and a majority of the non-partnered outliers would fall below the third quartile Q_3 of the partnered streamer group.

```
In [27]: cent_df = spark.createDataFrame(
    betweenness.items(), ["node", "centrality"])
cent_df = cent_df.join(
    other=att_df.select("new_id", "Partner"),
    on=[cent_df.node == att_df.new_id],
    how="inner",
)
plt.figure(figsize=(15, 9))
sns.boxplot(
    x=cent_df.select("centrality").rdd.map(lambda x: x[0]).collect(),
    hue=cent_df.select("partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("Betweenness Centrality By Partnership Status")
plt.xlabel("Betweenness Centrality Values")
plt.show()
```

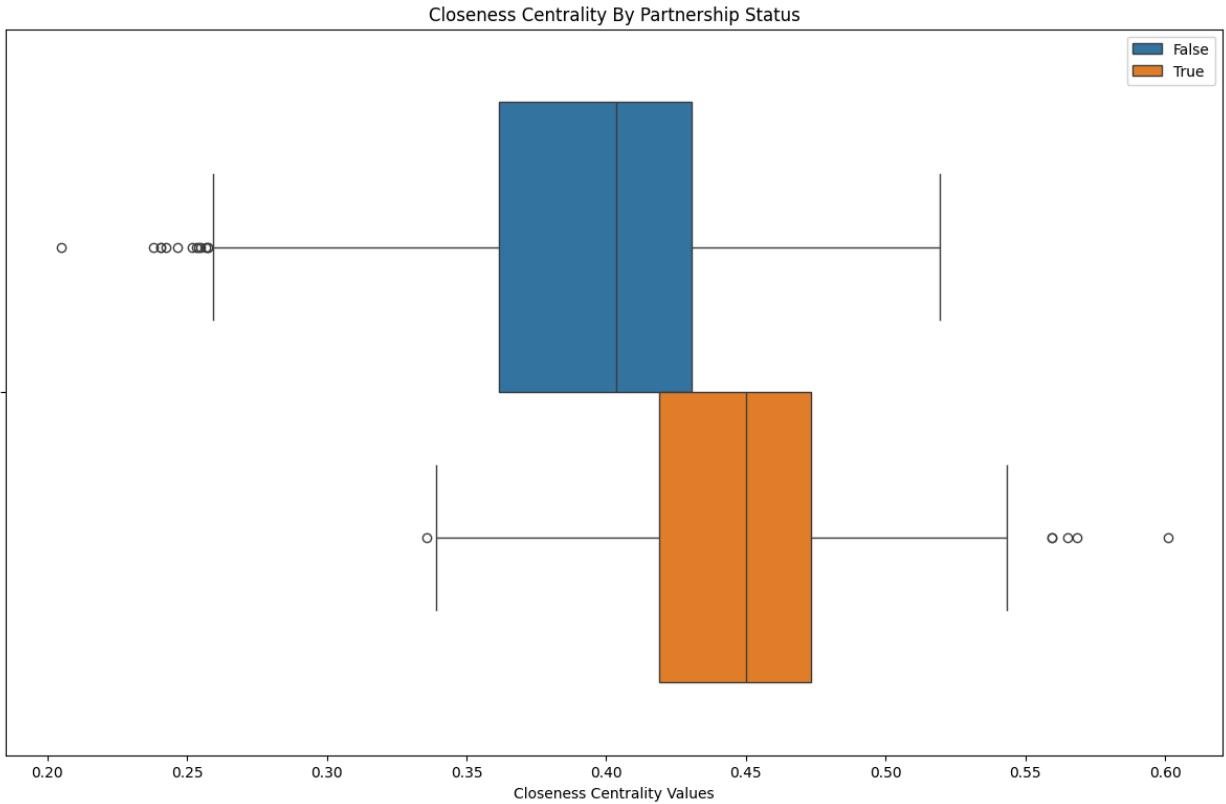
Betweenness Centrality By Partnership Status



Betweenness Centrality Boxplot

From the boxplot for *betweenness centrality* the partnered streamers are more important connecting to other user than the non-partnered streamers even though the betweenness in either group is not large.

```
In [28]: cent_df = spark.createDataFrame(
    closeness.items(), ["node", "centrality"])
cent_df = cent_df.join(
    other=att_df.select("new_id", "Partner"),
    on=[cent_df.node == att_df.new_id],
    how="inner",
)
plt.figure(figsize=(15, 9))
sns.boxplot(
    x=cent_df.select("centrality").rdd.map(lambda x: x[0]).collect(),
    hue=cent_df.select("partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("Closeness Centrality By Partnership Status")
plt.xlabel("Closeness Centrality Values")
plt.show()
```




Closeness Centrality Boxplot

The *closeness centrality* boxplot shows that partnered streamers have a greater ability to reach other users than the non-partnered streamers.

```
In [29]: cent_df = spark.createDataFrame(
    eigenvector.items(), ["node", "centrality"])
cent_df = cent_df.join(
    other=att_df.select("new_id", "Partner"),
    on=[cent_df.node == att_df.new_id],
    how="inner",
)
plt.figure(figsize=(15, 9))
sns.boxplot(
    x=cent_df.select("centrality").rdd.map(lambda x: x[0]).collect(),
    hue=cent_df.select("partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("Eigenvector Centrality By Partnership Status")
plt.xlabel("Eigenvector Centrality Values")
plt.show()
```


Eigenvector Centrality By Partnership Status



Eigenvector Centrality Boxplot

The *Eigenvector centrality* boxplot shows that partnered streamers are more influential than non-partnered streamers. The maximum and a majority of the non-partnered streamers outliers still fall below the third quartile $Q3$ of the partnered streamers showing us that partnered streamers are influential.

```
In [30]: cent_df = spark.createDataFrame(
    pagerank.items(), ["node", "centrality"])
cent_df = cent_df.join(
    other=att_df.select("new_id", "Partner"),
    on=[cent_df.node == att_df.new_id],
    how="inner",
)
plt.figure(figsize=(15, 9))
sns.boxplot(
    x=cent_df.select("centrality").rdd.map(lambda x: x[0]).collect(),
    hue=cent_df.select("partner").rdd.map(lambda x: x[0]).collect(),
)
plt.title("PageRank By Partnership Status")
plt.xlabel("PageRank Values")
plt.show()
```



PageRank Boxplot

The *PageRank* boxplot shows us that partnered streamers are more important to the friendships in the network than non-partnered streamers. Again, this is very heavily favored towards partnered streamers against non-partnered streamers.

Apply the Louvain Community Detection Algorithm Report the Number of Communities

After reading the documentation on NetworkX Louvain Community Algorithms:

- `louvain_communities`: returns the last iteration of the algorithms
- `louvain_partitions`: yields all iterations
 - Converting to a list results the same as `louvain_communities`

Note: A seed was set for both of the algorithms

```
In [31]: communities = nx.community.louvain_communities(G, seed=SEED)
print(f"Number of Communities: {len(communities)})")
```

Number of Communities: 6

```
In [32]: partitions = list(nx.community.louvain_partitions(G, seed=SEED))
print(f"Number of Partitions: {len(partitions)})")
print(f"Number of Partitions: {len(partitions[-1])})")
```

Number of Partitions: 3
Number of Partitions: 6

There are 6 communities from the Louvian community detection algorithm.

```
In [33]: print("Nodes Per Community")
for p, c in zip(partitions[-1], communities):
    print(f"Partition Algorithm:{len(p)}")
    print(f"Community Algorithm:{len(c)}")
    print()
```

```
Nodes Per Community
Partition Algorithm:96
Community Algorithm:96

Partition Algorithm:374
Community Algorithm:374

Partition Algorithm:485
Community Algorithm:485

Partition Algorithm:304
Community Algorithm:304

Partition Algorithm:437
Community Algorithm:437

Partition Algorithm:216
Community Algorithm:216
```

Compute the Average Views and Days per Community

Converting the list of sets to a dictionary with a community value.

```
In [34]: comm_dict = {}
for c in range(len(communities)):
    print(f"Current Node Count in Dictionary: {len(comm_dict)}")
    print(f"Adding Community {c + 1} Nodes: {len(communities[c])}")
    comm_dict.update(
        {(node, c + 1) for node in list(communities[c])})
print(f"New Node Count in Dictionary: {len(comm_dict)}\n")
```

```
Current Node Count in Dictionary: 0
Adding Comminity 1 Nodes: 96
New Node Count in Dictionary: 96

Current Node Count in Dictionary: 96
Adding Comminity 2 Nodes: 374
New Node Count in Dictionary: 470

Current Node Count in Dictionary: 470
Adding Comminity 3 Nodes: 485
New Node Count in Dictionary: 955

Current Node Count in Dictionary: 955
Adding Comminity 4 Nodes: 304
New Node Count in Dictionary: 1259
```

```
Current Node Count in Dictionary: 1259
Adding Comminity 5 Nodes: 437
New Node Count in Dictionary: 1696
```

```
Current Node Count in Dictionary: 1696
Adding Comminity 6 Nodes: 216
New Node Count in Dictionary: 1912
```

Creating a dataframe with communities and combining to attribute data

```
In [35]: comm_df = spark.createDataFrame(
    comm_dict.items(), ["node", "community"])

comm_df = comm_df.join(
    other=att_df.select("new_id", "Views", "Days"),
    on=[comm_df.node == att_df.new_id],
    how="inner",
)
comm_df.show(5)
```

node	community	new_id	Views	Days
80	1	80	9889	2284
34	1	34	6126	719
507	1	507	187209	809
586	1	586	3470	2332
1343	1	1343	294	1753

only showing top 5 rows

Create the dataframe with communities, average views, and average per day usage.

```
In [88]: window = Window.partitionBy()

comm_avg_df = comm_df.groupBy("community") \
    .agg(
        F.count("community").alias("community_count"),
        F.round(F.avg("Views"), 6).alias("avg_views"),
        F.sum("Days").alias("per_day"),
        F.round(F.avg("Days"), 6).alias("avg_per_days"),
    ) \
    .withColumn(
        "percent_avg_views",
        F.round(
```

```

        F.col("avg_views")
        / F.sum("avg_views").over(window), 4).cast(T.DoubleType())
        * 100
    ) \
    .withColumn(
        "percent_per_day",
        F.round(
            F.col("per_day")
            / F.sum("per_day").over(window), 4).cast(T.DoubleType())
            * 100
        ) \
        .withColumn(
            "percent_avg_days",
            F.round(
                F.col("avg_per_days")
                / F.sum("avg_per_days").over(window), 4).cast(T.DoubleType())
                * 100
            ) \
            .orderBy("community")
    comm_avg_df.show()

```

community	community_count	avg_views	per_day	avg_per_days	percent_avg_views	percent_per_day	percent_avg_days
1	96	293829.479167	166889	1738.427083	12.02	6.58	20.66999999999998
2	374	212264.254011	457668	1223.71123	8.68	19.9999999999	18.029
3	485	272537.934021	675075	1391.907216	11.15	26.6	14.54999999999999
4	304	1291529.848684	402944	1325.473684	52.82	99999999999	15.879
5	437	308141.036613	485237	1110.382151	12.6	19.12	13.20000000000001
6	216	66690.791667	350209	1621.337963	2.73	13.8	19.28

Average Views and Per Day Usage Per Community Observations

Community 1 with 96 user accounts is the smallest community with 12.02% of average views, but has the highest average per day with around 20.67%. This group accounts for 6.58% of days.

Community 2 with 374 users has one of the lowest average views with a third largest community at 8.68%, and has the second lowest per day with 14.56%. This group accounts for 18.03% of days.

Community 3 with 485 users is the largest community. The views and per day averages are in the middle of the pack with 11.15% and 16.55% respectively. This group accounts for 26.60% of days.

Community 4 with 304 users leads average views by a massive margin with 52.82% of all views, but has less average per day than communities 1, 3 and 6 with only around

15.88%. This group accounts for 18.03% of days.

Community 5 with 437 users is the second largest community with an average 12.60% of all views, but has the lowest average per day with only 13.20%. This group accounts for 19.12% of days.

Community 6 with 216 users is the second smallest community, and has the lowest average views with only 2.73% of all views, but has the second highest average per day with around 19.28%. This group accounts for 13.80% of days.

So from the table and information above, it would appear that communities 1 and 6 have the highest per day average in their community. Community 4 leads with the most views from a smaller user pool, and community 3 with the largest user group lead for per day usage, but only the third highest average per day usage.

End of Project